


None, null, nil: lessons from caching and representing nothing with something

Felice Ho
PyGotham 2019
New York City

I am Felice Ho



Overview

The scenario: what, why, where, how of caching


The problem with 'nothing'

'We have a problem'

Root cause analysis

Lessons

The value of 'nothing'



*“There are only two hard things
in Computer Science: cache
invalidation and naming things”*

- Phil Karlton

What happened was, **the new cat** went in for that little operation to ensure that he will be the Last of the Marlowes, and the vet offered us the option of either the ear-tattoo or implanted-microchip for permanent identification, recommending the microchip as more reliable (tattoos fade). This Microchip is I gather some sort of RFID technology, and as of now, Marlowe has a permanent unique identifier. I feel a new URI scheme coming on: just call little Marlowe **pet:cat:982009102637565**. My head is buzzing: Resource Description of Felines... POAF... cat semantics! The future awaits. [Update: It's not that easy; I should have known, as I've often quoted **Phil Karlton's** wise saying "There are only two hard things in Computer Science: cache invalidation and naming things". **Including pets**. (Thanks to Joe Pallas for the link.)]

"The first place anyone found it on the internet was in Tim Bray's blog.

Tim said that he first heard it around 1996-7"


- Martin Fowler

Source:
<https://martinfowler.com/bliki/TwoHardThings.html>

Goals of talk

How do we invalidate data in a cache?

How can production break down - from *nothing*!




What
is a cache and cache
invalidation?

Cache


Temporary storage of data
so future requests served
faster

Reduce server load,
improving app performance


Fetch data once, read more
than once



Buying milk at the supermarket



Server




Requests

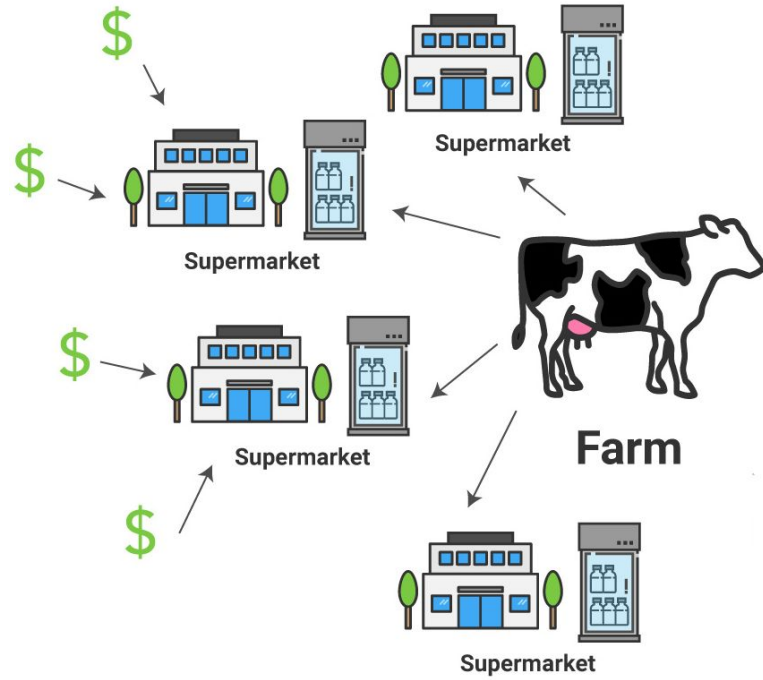
Buying milk at the supermarket

High demand
(large # of requests)

Limited supply of milk
(slow response, limited
computing power of
servers)



Supermarkets store milk



Reliable, quick to access

Expiration

Scalable

Markets (cache), not cows
handle demand of
consumers

Cache invalidation

Data changes without you knowing about it

Whether there is a change in data, no data, or new data

Cache needs to get updated


Cache invalidation

Data in cache is temporary

Cache needs to get updated or
data removed

Market needs to know when
milk is expired (not valid) and
to remove from shelves
(delete from cache)






Why
are we caching and how can
it help you with app
performance?

The problem

Slow website and app performance

Multiple data sources

Single request requires data from different systems




The problem

Multiple applications

- accessing same data content in different ways
- running the same queries at different times

High burden/load on databases



The ask

Build a cache for quick retrieval of data

Make it easier to build high performing web applications with fewer errors and quicker response times

The ask

Relieve SQL load on databases

Easier and more reliable path to data

Consolidated data, source of truth, consistent data across all applications



Where
are we caching data and how
can it be accessed?

Cache Storage - Redis


Open source, in memory data structure store

Built-in replication


Highly available

Fault tolerant


Highly scalable



Cache Access - API - fetch data



Cache Access - API - cache miss





How
are we caching data?

Factors to consider - API contract

Agreement between service and client

Specifications on data and structure

-> need a JSON response string



The screenshot shows a web browser's developer tools interface. The 'Body' tab is selected, displaying a JSON object. The JSON is formatted in a 'Pretty' view, showing a root object with a 'profile' property. The 'profile' property is an object containing 'city', 'conference', and 'name' properties. The 'city' is 'new york city', the 'conference' is 'pygotham 2019', and the 'name' is 'felice ho'. The interface includes tabs for 'Body', 'Cookies', 'Headers (4)', and 'Test Results'. Below the tabs are buttons for 'Pretty', 'Raw', and 'Preview', along with a 'JSON' dropdown and a refresh icon.

```
{
  "profile": {
    "city": "new york city",
    "conference": "pygotham 2019",
    "name": "felice ho"
  }
}
```


Factors to consider - API design

Strict structure (include null values)

- Indicate existence of property even if there is no value
- Would need to handle non-nullable fields

```
{  
  "profile": {  
    "city": "new york city",  
    "conference": "pygotham 2019",  
    "name": "felice ho",  
    "prior_talks": null  
  }  
}
```

Factors to consider - API design

Dynamic structure (do not include null values)

- Removes noise, omission generally represents lack of value
- Unclear if omissions mean unknown or truly no value

```
{  
  "profile": {  
    "city": "new york city",  
    "conference": "pygotham 2019",  
    "name": "felice ho"  
  }  
}
```

Factors to consider - API design

Dynamic or strict structure

Also depends on data - sparse or dense data

-> null values not included in API response

Factors to consider - the data itself

Transactional / point of sale

Web and application data

CRM

Data warehouse

-> update strategy needs to
include all data sources



Factors to consider - update strategy

Cache warming - daily batch cache updates

Time to live (TTL)

- proactively assign expiration
- avoid unexpected *LRU eviction of data

* Least Recently Used

Factors to consider - update strategy

Cache miss functionality in API

-> need to ensure accurate and relevant data in cache

Caching options in Redis

```
# Python interface to the Redis key-value store.  
# pip install redis  
import redis  
redis_conn = redis.Redis(  
    host='localhost', port=6379, db=0  
)
```

Caching options in Redis

```
# Cache as json encoded string
import json

# hset pythonista.1234567890 profile <json_str>
profile_dict = {
    "name": "felice",
    "city": "new york city",
    "conference": "pygotham 2019",
    "prior_talks": None
}
profile_json = json.dumps(profile_dict)

redis_conn.hset(
    name='pythonista.1234567890',
    key='profile',
    value=profile_json
)
```

```
127.0.0.1:6379> hgetall pythonista.
1234567890
1) "profile"
2) "{\"name\": \"felice\", \"city\": \"new york city\", \"conference\": \"pygotham 2019\", \"prior_talks\": null}"
```

Redis

Caching options in Redis

```
# Cache via ReJSON: native JSON data type
# pip install rejson
from rejson import Client, Path

rj_conn = Client(
    host='localhost',
    port=6379,
    decode_responses=True
)

# JSON.SET pythonista.1234567890 . '{"profile": {
profile_dict = {
    "name": "felice",
    "city": "new york city",
    "conference": "pygotham 2019",
    "prior_talks": None
}
obj = {"profile": profile_dict}

rj_conn.jsonset(
    'pythonista.1234567890', Path.rootPath(), obj
)
```

```
127.0.0.1:6379> JSON.GET pythonista
.1234567890
"{\"profile\":{\"name\":\"felice\",
\"city\":\"new york city\", \"confer
ence\":\"pygotham 2019\", \"prior_ta
lks\":null}}"
```

Redis

Caching options in Redis

```
# Cache as hashes, Redis key/value pairs
# hmset pythonista.1234567890:profile name
profile_dict = {
    "name": "felice",
    "city": "new york city",
    "conference": "pygotham 2019",
    "prior_talks": "null"
}

redis_conn.hmset(
    name='pythonista.1234567890:profile',
    mapping=profile_dict
)
```

```
127.0.0.1:6379> hgetall pythonista.
1234567890:profile
1) "name"
2) "felice"
3) "city"
4) "new york city"
5) "conference"
6) "pygotham 2019"
7) "prior_talks"
8) "null"
```

Redis

Note: Starting with redis-py 3.0, `None` is no longer accepted as input for keys or values. Same for `True` or `False`. Users will need to cast these values explicitly before sending them to redis-py.

Source: <https://github.com/andymccurdy/redis-py/issues/190>

Caching strategy

Tried both JSON string and hashes in production

No notable performance difference

- hashes slightly faster with help of Lua and cJSON

The problem with 'nothing'

Representing nothing with something

Keep placeholder value for keys even if null

Recognize data changed from existing to not existing

Else appears as if something exists, when it doesn't causing invalid data in cache

What is null?

Value assigned to a variable to represent

- No value, neutral behavior, absence of data, non value, absence of useful value, nothing

What is null?

Represented with zeros
but not same value as
zero

ASCII control characters			ASCII printable characters					
00	NULL	(Null character)	32	space	64	@	96	`
01	SOH	(Start of Header)	33	!	65	A	97	a
02	STX	(Start of Text)	34	"	66	B	98	b
03	ETX	(End of Text)	35	#	67	C	99	c
04	EOT	(End of Trans.)	36	\$	68	D	100	d
05	ENQ	(Enquiry)	37	%	69	E	101	e
06	ACK	(Acknowledgement)	38	&	70	F	102	f
07	BEL	(Bell)	39	'	71	G	103	g
08	BS	(Backspace)	40	(72	H	104	h
09	HT	(Horizontal Tab)	41)	73	I	105	i
10	LF	(Line feed)	42	*	74	J	106	j
11	VT	(Vertical Tab)	43	+	75	K	107	k
12	FF	(Form feed)	44	,	76	L	108	l
13	CR	(Carriage return)	45	-	77	M	109	m
14	SO	(Shift Out)	46	.	78	N	110	n
15	SI	(Shift In)	47	/	79	O	111	o
16	DLE	(Data link escape)	48	0	80	P	112	p
17	DC1	(Device control 1)	49	1	81	Q	113	q

What is the problem here?

Serialization of null is represented differently

Python: The sole value of the type `NoneType`. `None` is frequently used to represent the absence of a value as when default arguments are not passed to a function. Assignments to `None` are illegal and raise a `SyntaxError`.

What is the problem here?

Serialization of null is represented differently

JSON:

A value can be a string in double quotes, or a number, or true or false or null, or an object or an array. These structures can be nested.

What is the problem here?

Serialization of null is represented differently

Lua: Nil is a type with a single value, **nil**, whose main property is to be different from any other value. As we have seen, a global variable has a **nil** value by default, before a first assignment, and you can assign **nil** to a global variable to delete it. Lua uses **nil** as a kind of non-value, to represent the absence of a useful value.

What is the problem here?


Serialization of null is represented differently

Redis: it's not possible!

Redis treats everything as a string. **It has no concept of Null values.** Even when fetching a key that doesn't exist, a sane default for that key type is returned rather than a Null value. For instance, Redis specifies that a LRANGE command against a key that doesn't exist returns an empty list.

What is the problem here?

It is up to the language or library to determine how to represent null



Storing null in Redis

Store as empty string ""

- Is value actually an empty string or null



""



0



null

undefined

Storing null in Redis

Use sentinel value to represent null

- “null” not same as null




Photo credit: <https://xkcd.com/327/>

Trust your encoder

Encoder will serialize into what you need

Keep null values in native format before encoding

Python	Redis	Lua	JSON
None	(your choice)	nil	null

Trust your encoder

Mapping key/value pairs from Python dictionary to JSON string

```
import json

# Python to JSON
profile_dict = {
    "name": "felice",
    "city": "new york city",
    "conference": "pygotham 2019",
    "prior_talks": None
}

profile_json = json.dumps(profile_dict)
print(profile_json)
```

```
{"name": "felice", "city": "new york city",
"conference": "pygotham 2019", "prior_talks": null}
```


Trust your encoder

Mapping key/pair values
from Redis using Lua to
encode to JSON string

```
-- Redis to Lua
local keyvalues = redis.call('HGETALL', keyname_category);
local category_result = {};

for k = 1, #keyvalues, 2 do
    local key = keyvalues[k]
    local value = keyvalues[k + 1]

    if value == <your-redis-sentinel-value> then
        category_result[key] = nil
        -- to show nulls use cJSON.null as value
    else
        category_result[key] = value
    end
end

return cJSON.encode(category_result);
```

What is serialization and encoding?

Encoder

- Helps with converting data into a certain representation, from one format into another (Python to JSON)




Photo credit: <http://bit.ly/2oAEdFh>

What is serialization and encoding?

Serialization

- Process of translating an object into a format that can be stored or transmitted, and reconstructed later (JSON is a serialization format for client server communication)



‘We have a problem’

◇ Message

Could not convert string to DateTime: null. Path 'prospectExpirationDate', line 1, position 113.

Could not convert string to DateTime: null. Path 'prospectExpirationDate', line 1, position 110.

Could not convert string to DateTime: null. Path 'prospectExpirationDate', line 1, position 110.

Could not convert string to DateTime: null. Path 'prospectExpirationDate', line 1, position 111.

Storing null in Redis

Redis treats everything as a string. **It has no concept of Null values.** Even when fetching a key that doesn't exist, a sane default for that key type is returned rather than a Null value. For instance, Redis specifies that a LRANGE command against a key that doesn't exist returns an empty list.

“null” not same as null

Root cause analysis



Lessons

Represent nothing with something

Nothing is recognized differently

Handle non-nullable values appropriately

Be aware of how your source data systems and tools handle null values

Left join (important) data

Include records with values at one point,

- that matter if they no longer have values,
- or otherwise not removed via TTL

Trust your encoder

Serialize 'nothing' in native form, no matter which language, tool, or format you are using



The value of 'nothing'

"I call it my billion-dollar mistake. It was the invention of the null reference in 1965. ... I couldn't resist the temptation to put in a null reference, simply because it was so easy to implement. This has led to innumerable errors, vulnerabilities, and system crashes, which have probably caused a billion dollars of pain and damage in the last forty years."

- Tony Hoare
2009

Embracing null

Useful for cache invalidation

Web applications

- Reduced errors, quicker response times

Databases

- Reduced SQL load



Thank you!