


None, null, nil: lessons from caching & representing nothing with something

Felice Ho
PyTennessee 2020
Nashville, TN

I am Felice Ho



Overview

The scenario: what, why, where, how of caching


The problem with 'nothing'

'We have a problem'

Root cause analysis

Lessons

The value of 'nothing'



*“There are only two hard things
in Computer Science: cache
invalidation and naming things”*

- Phil Karlton

What happened was, **the new cat** went in for that little operation to ensure that he will be the Last of the Marlowes, and the vet offered us the option of either the ear-tattoo or implanted-microchip for permanent identification, recommending the microchip as more reliable (tattoos fade). This Microchip is I gather some sort of RFID technology, and as of now, Marlowe has a permanent unique identifier. I feel a new URI scheme coming on: just call little Marlowe **pet:cat:982009102637565**. My head is buzzing: **Resource Description of Felines... POAF... cat semantics!** The future awaits. *[Update: It's not that easy; I should have known, as I've often quoted **Phil Karlton's** wise saying "There are only two hard things in Computer Science: cache invalidation and naming things". Including pets. (Thanks to Joe Pallas for the link.)]*

Comments on this fragment are closed.

Updated: 2005/12/23

ongoing

What this is · 

Truth · **Biz** · **Tech**

author · **Dad** · **software** ·
colophon · **rights**



December 23, 2005

- **The World** (116 fragments)
- • **Humor** (23 more)
- **Technology** (85 fragments)
- • **Web** (390 more)

By **Tim Bray**

I am an employee of Amazon.com, but the opinions expressed here are my own, and no other party necessarily agrees with them.

A full disclosure of my professional interests is on the **author** page.

*“The first place anyone
found it on the internet was
in Tim Bray's blog.*


*Tim said that he first heard it
around 1996-7”*

- Martin Fowler


Goals of talk


How do we invalidate data in a cache?

How can production break down - from *nothing*!



What
is a cache and cache
invalidation?






Application




Database


Reduce server load, improve app performance




Fetch data once, read more than once




Buying milk at the supermarket



Server



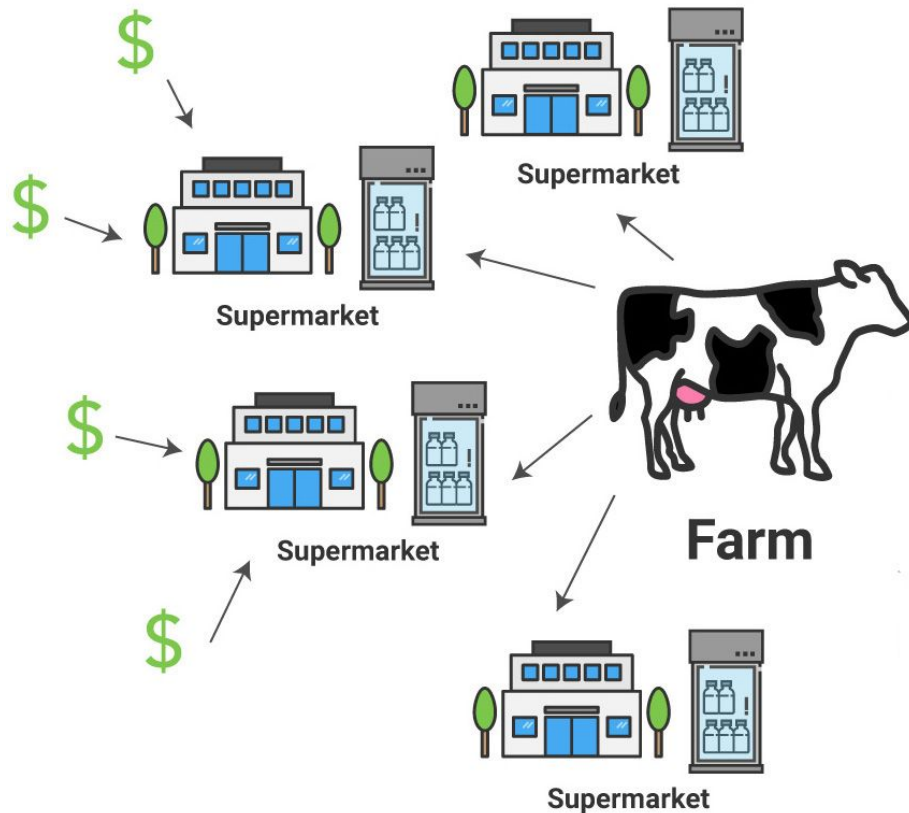
Requests



High demand

Limited supply of milk

Supermarkets store milk



Reliable, quick to access

Expiration

Scalable

Markets handle demand
of consumers

Cache invalidation

Data changes without you knowing about it

Whether there is a change in data, no data, or new data

Cache needs to get updated


Cache invalidation

Data in cache is temporary

Cache needs to get updated or
data removed

Market needs to know when
milk is expired and to remove
from shelves






Why
are we caching and how can
it help you with app
performance?

The problem

Slow website and app performance

- multiple data sources

Single request requires data from different systems




The problem

Multiple web applications

- accessing exact same data in different ways
- running similar queries at different times

High burden / load on databases



The problem

Third-party provider APIs

- rate limits
- slowness
- token issues

Errors in business systems and digital products

The ask

Build a cache for quick retrieval of data

Make it easier to build high performing web applications with fewer errors and quicker response times

The ask

Relieve SQL load on databases

Easier and more reliable path to data

Consolidated data, source of truth, consistent data across all applications



Where
are we caching data and how
can it be accessed?

Cache storage - Redis


Open source, in memory data structure store

Built-in replication


Highly available

Fault tolerant


Highly scalable



Cache access - API - fetch data



Cache access - API - cache miss





How
are we caching data?

Factors to consider - API contract

Agreement between service and client

Specifications on data and structure

-> need a JSON response string



The screenshot shows a web browser's developer tools interface. The 'Body' tab is selected, displaying a JSON response. The JSON is formatted in a 'Pretty' view, showing a root object with a 'profile' property. The 'profile' object contains three properties: 'city' (nashville), 'conference' (pytennessee 2020), and 'name' (felice). The interface includes tabs for 'Body', 'Cookies', 'Headers (4)', and 'Test Results'. Below the tabs are buttons for 'Pretty', 'Raw', and 'Preview', along with a 'JSON' dropdown menu and a refresh icon. A line number indicator on the left shows lines 1 through 7.

```
1 {  
2   "profile": {  
3     "city": "nashville",  
4     "conference": "pytennessee 2020",  
5     "name": "felice"  
6   }  
7 }
```

Factors to consider - API design

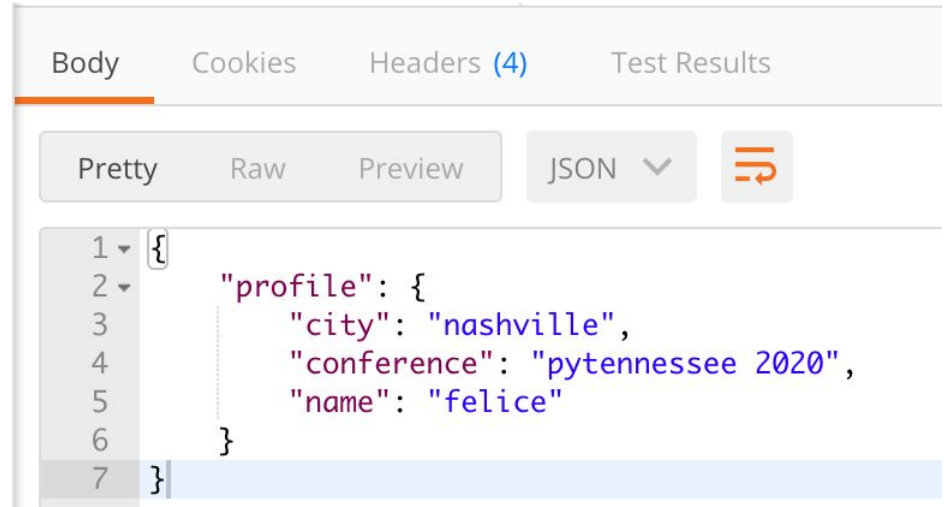
Dynamic or strict structure

- to not include or include null values

Factors to consider - API design

Dynamic structure

- removes noise, omission represents lack of value
- unclear if omissions mean unknown or truly no value



```
1 {  
2   "profile": {  
3     "city": "nashville",  
4     "conference": "pytennessee 2020",  
5     "name": "felice"  
6   }  
7 }
```

Factors to consider - API design

Strict structure

- indicates existence of property even if there is no value
- would need to handle non-nullable fields



```
{
  "profile": {
    "city": "nashville",
    "conference": "pytennessee 2020",
    "name": "felice",
    "prior_talks": "null"
  }
}
```


Factors to consider - API design

Dynamic or strict structure

It depends...

sparse or dense data?

-> null values not included in API response

Factors to consider - the data itself

Transactional / point of sale

Web and application data

CRM

Data warehouse

-> update strategy needs to
include all data sources



Factors to consider - update strategy

Cache warming

Time to live (TTL)

Cache miss functionality in API

-> need to ensure accurate and relevant data in cache

Caching options in Redis

```
# Python interface to the Redis key-value store.  
# pip install redis  
import redis  
redis_conn = redis.Redis(  
    host='localhost', port=6379, db=0  
)
```

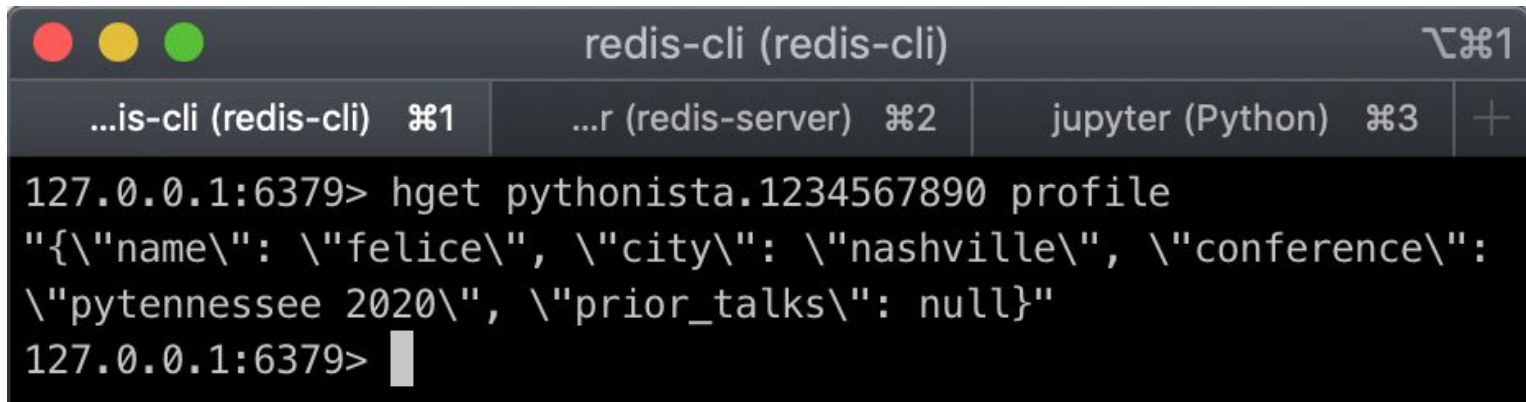
Caching options in Redis

```
# Cache as json encoded string
import json

# hset pythonista.1234567890 profile '{"name": "felice ho", "city": "nashv.
profile_dict = {
    "name": "felice",
    "city": "nashville",
    "conference": "pytennessee 2020",
    "prior_talks": None
}
profile_json = json.dumps(profile_dict)

redis_conn.hset(
    name='pythonista.1234567890',
    key='profile',
    value=profile_json
)
```

Caching options in Redis



```
redis-cli (redis-cli) 127.0.0.1:6379> hget pythonista.1234567890 profile
{"name": "felice", "city": "nashville", "conference": "pytennessee 2020", "prior_talks": null}
```

The screenshot shows a macOS-style window titled "redis-cli (redis-cli)" with three colored window control buttons (red, yellow, green) in the top-left corner. The window has a dark gray title bar and a black terminal area. The terminal shows a Redis command prompt "127.0.0.1:6379>" followed by the command "hget pythonista.1234567890 profile". The output is a JSON string: {"name": "felice", "city": "nashville", "conference": "pytennessee 2020", "prior_talks": null}. The prompt "127.0.0.1:6379>" is repeated on the next line with a white cursor. Above the terminal area, there is a tab bar with three tabs: "...is-cli (redis-cli) 1", "...r (redis-server) 2", and "jupyter (Python) 3". A plus sign icon is to the right of the tabs.

Caching options in Redis

```
# Cache via ReJSON: native JSON data type
# pip install rejson
from rejson import Client, Path

rj_conn = Client(
    host='localhost',
    port=6379, decode_responses=True
)

# JSON.SET pythonista.1234567890 . '{"profile": {"name": "felice ho", "city": "nashville", "conference": "pytennessee 2020", "prior_talks": None}}'
profile_dict = {
    "name": "felice",
    "city": "nashville",
    "conference": "pytennessee 2020",
    "prior_talks": None
}
obj = {"profile": profile_dict}

rj_conn.jsonset(
    'pythonista.1234567890', Path.rootPath(), obj
)
```

Caching options in Redis



```
redis-cli (redis-cli) 127.0.0.1:6379> JSON.GET pythonista.1234567890
{"profile":{"name":"felice","city":"nashville","conference":"pytennessee 2020","prior_talks":null}}
```

The screenshot shows a macOS-style terminal window titled "redis-cli (redis-cli)". The terminal has three tabs: "redis-cli (redis-... 1", "docker (docker) 2", and "jupyter (Pyt... 3". The active tab is "redis-cli (redis-... 1". The terminal content shows a Redis command prompt "127.0.0.1:6379>" followed by the command "JSON.GET pythonista.1234567890". The output is a JSON string: "{\"profile\":{\"name\":\"felice\",\"city\":\"nashville\",\"conference\":\"pytennessee 2020\",\"prior_talks\":null}}\"". The prompt "127.0.0.1:6379>" is followed by a cursor.

Caching options in Redis

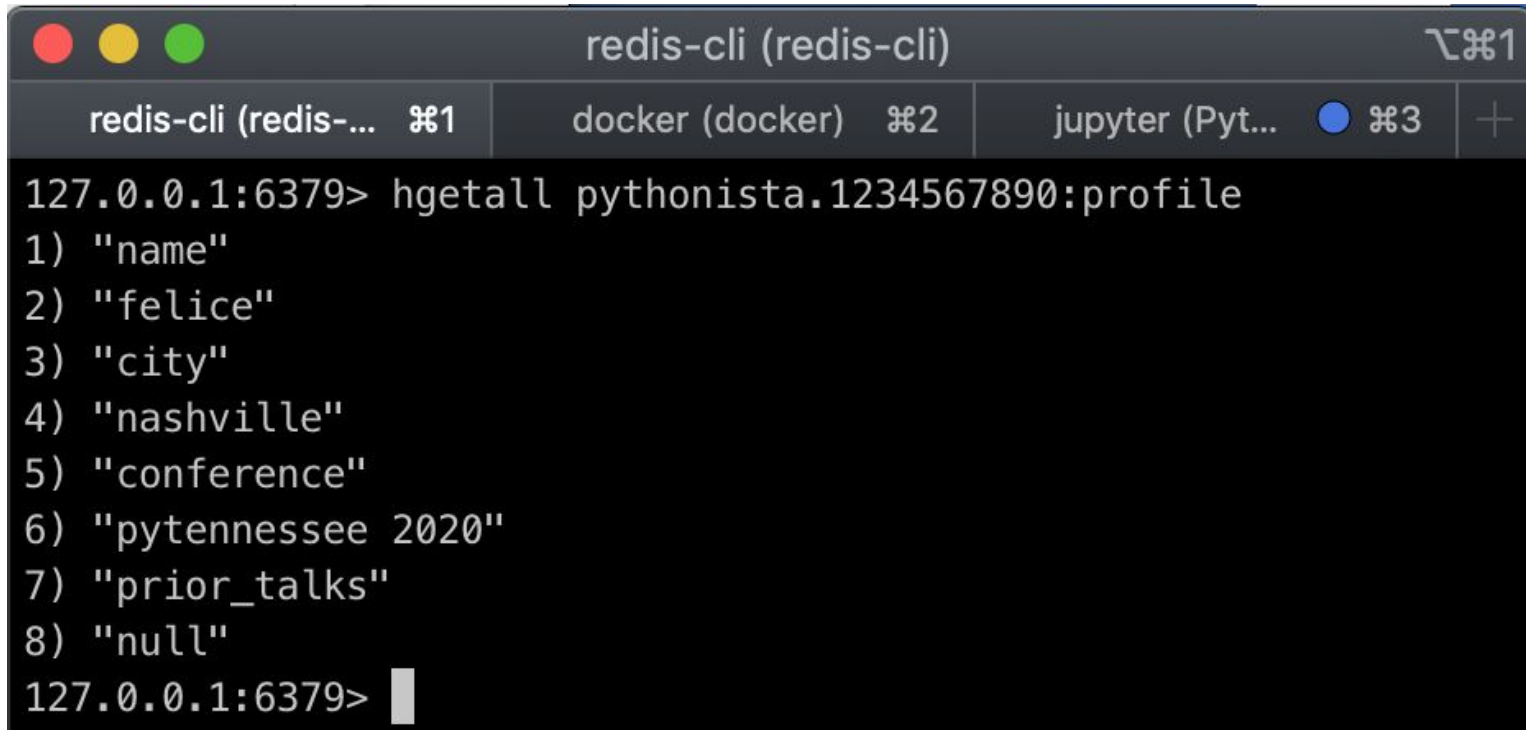
```
# Cache as hashes, Redis key/value pairs
# hmset pythonista.1234567890:profile name 'felice ho' city 'nashville' co.
profile_dict = {
    "name": "felice",
    "city": "nashville",
    "conference": "pytennessee 2020",
    "prior_talks": "null"
}

redis_conn.hmset(
    name='pythonista.1234567890:profile',
    mapping=profile_dict
)
```

Note: Starting with redis-py 3.0, `None` is no longer accepted as input for keys or values. Same for `True` or `False`. Users will need to cast these values explicitly before sending them to redis-py.

Source: <https://github.com/andymccurdy/redis-py/issues/190>

Caching options in Redis



```
redis-cli (redis-cli) 127.0.0.1:6379> hgetall pythonista.1234567890:profile
1) "name"
2) "felice"
3) "city"
4) "nashville"
5) "conference"
6) "pytennessee 2020"
7) "prior_talks"
8) "null"
127.0.0.1:6379>
```

The image shows a terminal window titled "redis-cli (redis-cli)". The terminal displays the command `hgetall pythonista.1234567890:profile` and its output, which is a list of 8 items: `1) "name"`, `2) "felice"`, `3) "city"`, `4) "nashville"`, `5) "conference"`, `6) "pytennessee 2020"`, `7) "prior_talks"`, and `8) "null"`. The prompt `127.0.0.1:6379>` is visible at the bottom.

Caching strategy

JSON string vs. hashes

- no notable performance difference
- hashes slightly faster with help of Lua and cJSON

The problem with 'nothing'

Representing nothing with something

Keep placeholder value for keys even if null

Recognize data changed from existing to not existing

Else appears as if something exists, when it doesn't causing invalid data in cache

What is null?

Value assigned to a variable to represent

- no value / non value
- neutral behavior
- absence of data / useful value
- nothing

What is null?

Represented with zeros
but not same value as
zero

ASCII control characters			ASCII printable characters					
00	NULL	(Null character)	32	space	64	@	96	`
01	SOH	(Start of Header)	33	!	65	A	97	a
02	STX	(Start of Text)	34	"	66	B	98	b
03	ETX	(End of Text)	35	#	67	C	99	c
04	EOT	(End of Trans.)	36	\$	68	D	100	d
05	ENQ	(Enquiry)	37	%	69	E	101	e
06	ACK	(Acknowledgement)	38	&	70	F	102	f
07	BEL	(Bell)	39	'	71	G	103	g
08	BS	(Backspace)	40	(72	H	104	h
09	HT	(Horizontal Tab)	41)	73	I	105	i
10	LF	(Line feed)	42	*	74	J	106	j
11	VT	(Vertical Tab)	43	+	75	K	107	k
12	FF	(Form feed)	44	,	76	L	108	l
13	CR	(Carriage return)	45	-	77	M	109	m
14	SO	(Shift Out)	46	.	78	N	110	n
15	SI	(Shift In)	47	/	79	O	111	o
16	DLE	(Data link escape)	48	0	80	P	112	p
17	DC1	(Device control 1)	49	1	81	Q	113	q

What is the problem here?

Serialization of null is represented differently

Python:

The sole value of the type `NoneType`. `None` is frequently used to represent the absence of a value, as when default arguments are not passed to a function. Assignments to `None` are illegal and raise a `SyntaxError`.

What is the problem here?

Serialization of null is represented differently

JSON:

A value can be a *string* in double quotes, or a *number*, or **true** or **false** or **null**, or an *object* or an *array*. These structures can be nested.

What is the problem here?

Serialization of null is represented differently

Lua:

Nil is a type with a single value, **nil**, whose main property is to be different from any other value. As we have seen, a global variable has a **nil** value by default, before a first assignment, and you can assign **nil** to a global variable to delete it. Lua uses **nil** as a kind of non-value, to represent the absence of a useful value.

What is the problem here?


Serialization of null is represented differently

Redis: it's not possible!

Redis treats everything as a string. It has no concept of Null values. Even when fetching a key that doesn't exist, a sane default for that key type is returned rather than a Null value. For instance, Redis specifies that a LRANGE command against a key that doesn't exist returns an empty list.

What is the problem here?

It is up to the language or library to determine how to represent null



Storing null in Redis

Store as empty string

- is value actually an empty string or null



""



0



null

undefined

Storing null in Redis

Use sentinel value to represent null

- “null” not same as null



Photo credit: <https://xkcd.com/327/>

Trust your encoder

Encoder will serialize into what you need

Keep null values in native format before encoding




None



null



nil



(your choice)

Trust your encoder

Mapping key/value pairs from Python dictionary to JSON string

```
import json

# Python to JSON
profile_dict = {
    "name": "felice",
    "city": "nashville",
    "conference": "pytenessee 2020",
    "prior_talks": None
}
profile_json = json.dumps(profile_dict)
print(profile_json)
```

```
{"name": "felice", "city": "nashville",
 "conference": "pytenessee 2020", "prior_talks": null}
```


Trust your encoder

Mapping key/pair values
from Redis using Lua to
encode to JSON string

```
-- Redis to Lua
local keyvalues = redis.call('HGETALL', keyname_category);
local category_result = {};

for k = 1, #keyvalues, 2 do
    local key = keyvalues[k]
    local value = keyvalues[k + 1]

    if value == <your-redis-sentinel-value> then
        category_result[key] = nil
        -- to show nulls use cJSON.null as value
    else
        category_result[key] = value
    end
end

return cJSON.encode(category_result);
```

Encoding vs. Serialization

Encoder

Helps with converting data into a certain representation,
from one format into another

-> Python to JSON

Serialization

Process of translating an object into a format that can be stored or transmitted, and reconstructed later

-> JSON is a serialization format for client server communication



‘We have a problem’

Message

Could not convert string to DateTime: null. Path 'prospectExpirationDate', line 1, position 113.

Could not convert string to DateTime: null. Path 'prospectExpirationDate', line 1, position 110.

Could not convert string to DateTime: null. Path 'prospectExpirationDate', line 1, position 110.

Could not convert string to DateTime: null. Path 'prospectExpirationDate', line 1, position 111.

Storing null in Redis

Redis treats everything as a string. It has no concept of Null values. Even when fetching a key that doesn't exist, a sane default for that key type is returned rather than a Null value. For instance, Redis specifies that a LRANGE command against a key that doesn't exist returns an empty list.

“null” not same as null

Root cause analysis



Photo credit: Abbott and Costello

Lessons

Represent nothing with something

Nothing is recognized differently

Handle non-nullable values appropriately

Be aware of how your source data systems and tools handle null values

Left join (important) data

Include records

- with values at one point
- that matter if they no longer have values
- or otherwise not removed via TTL

Trust your encoder

Serialize 'nothing' in native form, no matter which language, tool, or format you are using



The value of 'nothing'

"I call it my billion-dollar mistake. It was the invention of the null reference in 1965. ... I couldn't resist the temptation to put in a null reference, simply because it was so easy to implement. This has led to innumerable errors, vulnerabilities, and system crashes, which have probably caused a billion dollars of pain and damage in the last forty years."

- Tony Hoare
2009

Embracing null

Useful for cache invalidation

Web applications

- reduced errors, quicker response times

Databases

- reduced SQL load

Thank you!

Slides:

<http://bit.ly/pytn-nonnullnil>

Blog post:

<http://bit.ly/globey-part-one>

LinkedIn:

 [/feliceho/](https://www.linkedin.com/company/feliceho/)